

Hydra Operator Guide

Table of Contents

Preface	1
1. Part I: Understanding Hydra for Data Operations	2
2. Operator Fundamentals	3
2.1. What is a Hydra Operator?	3
2.1.1. Operator Responsibilities	3
2.1.2. Typical Use Cases	3
2.2. Core Concepts for Operators	4
2.2.1. The Hydra Data Model	4
2.2.2. Schemas and Constraints	4
2.2.3. Mappings and Transformations	5
2.2.4. Data Quality and Validation	6
2.3. Development Environment and Tools	7
2.3.1. Setting Up Your Environment	7
2.3.2. Development Workflow	8
2.3.3. Available Tools and Resources	8
2.4. Project Organization	9
2.4.1. Recommended Directory Structure	9
2.4.2. Configuration Management	10
2.5. Best Practices	11
2.5.1. Schema Design	11
2.5.2. Transformation Design	11
2.5.3. Performance Considerations	12
2.6. Common Patterns and Anti-Patterns	12
2.6.1. Effective Patterns	12
2.6.2. Anti-Patterns to Avoid	13
2.7. Next Steps	13
2.8. Related Documentation	13
3. Part II: Getting Started with Data Transformation	14
4. Transformation Examples	15
4.1. Your First Transformation: Relational to Graph	15
4.1.1. The Scenario	15
4.1.2. Step 1: Define Your Input Schema	15
4.1.3. Step 2: Define Your Output Schema	16
4.1.4. Step 3: Define Your Mapping	17
4.1.5. Step 4: Generate Your Graph	18
4.2. Advanced Transformation Patterns	19
4.2.1. Working with Complex Data Types	19
4.2.2. Multi-Step Transformations	19

4.2.3. Schema Evolution and Migration	21
4.3. Working with Different Data Formats	21
4.3.1. CSV and Tabular Data	21
4.3.2. JSON and Semi-Structured Data	22
4.3.3. XML and Hierarchical Data	23
4.4. Message Translation and Integration	23
4.4.1. Building Message Translators	23
4.4.2. Event Stream Processing	24
4.4.3. Data Validation and Quality	25
4.5. Performance and Optimization	26
4.5.1. Streaming and Large Datasets	26
4.5.2. Caching and Memoization	27
4.6. Related Documentation	28
5. Part III: Development Workflow with Pixi Tasks	29
5.1. Working with Generated Code	29
5.2. Data Processing Tasks	29
5.3. Testing Your Transformations	29
5.4. Exploring Hydra Modules and Functions	30
5.5. Building and Packaging	30
5.6. Advanced Workflow Patterns	30
5.6.1. Iterative Development and Testing	30
5.6.2. Working with Large Datasets	30
5.6.3. Debugging Transformations	31
5.6.4. Performance Optimization	31
6. Part IV: Production Deployment	32
6.1. Packaging for Production	32
6.1.1. Creating Distributable Packages	32
6.1.2. Configuration Management	32
6.2. Monitoring and Maintenance	32
6.2.1. Health Monitoring	32
6.2.2. Maintenance Tasks	32
7. Appendices	33
Appendix A: Operator Resources	33
Appendix B: Common Transformation Patterns	33
7.B.1. Data Format Conversions	33
7.B.2. Schema Evolution Patterns	33
7.B.3. Performance Optimization Patterns	33
Appendix C: Troubleshooting Guide	33
7.C.1. Common Issues	33
7.C.2. Getting Help	34
Colophon	35

Preface

This guide is for developers who are using Hydra to create message translators using coders and mappers. If you're working on the core Hydra kernel, see the [Developer Guide](#). If you're creating new Hydra language implementations, see the [Implementor Guide](#).

Hydra empowers users with a toolkit of functions and DSLs (Domain Specific Languages) to transform and validate graph data and schemas with maximum versatility across types and standards. In Josh Shinavier's words: "In Hydra, programs are graphs, and graphs are programs."

This guide will walk you through practical examples of using Hydra for data transformation tasks, from simple relational-to-graph conversions to complex schema mappings and message translation scenarios.

Chapter 1. Part I: Understanding Hydra for Data Operations

Chapter 2. Operator Fundamentals

This document covers the fundamental concepts and core principles for operators using Hydra to create data transformation pipelines and message translators.

2.1. What is a Hydra Operator?

A Hydra operator is someone who uses existing Hydra implementations to solve practical data transformation problems. Unlike core developers (who work on Hydra itself) or implementors (who create new language bindings), operators focus on applying Hydra's capabilities to real-world data integration challenges.

2.1.1. Operator Responsibilities

- **Schema Design:** Define input and output data structures
- **Mapping Creation:** Specify how data flows between different formats
- **Pipeline Implementation:** Build complete data transformation workflows
- **Quality Assurance:** Ensure data integrity and transformation correctness
- **Performance Optimization:** Optimize transformations for production workloads
- **Monitoring & Maintenance:** Monitor running systems and handle issues

2.1.2. Typical Use Cases

Data Integration

- Converting between different data formats (CSV ↔ JSON ↔ XML)
- Migrating data between systems with different schemas
- Aggregating data from multiple sources
- Real-time data transformation in streaming pipelines

Message Translation

- API format conversion (REST ↔ GraphQL ↔ Protocol Buffers)
- Legacy system integration
- Microservice communication adaptation
- Event format standardization

Schema Evolution


- Database migration scripts
- API versioning and compatibility
- Backward/forward compatibility handling

- Schema validation and enforcement

2.2. Core Concepts for Operators

2.2.1. The Hydra Data Model

Hydra uses a unified core data model that serves as an intermediate representation for all data transformations. Understanding this model is crucial for effective operations.



Type System

- **Primitive Types:** String, Integer, Float, Boolean, Binary
- **Composite Types:** Records (structs), Unions (sum types), Lists, Maps, Sets
- **Optional Types:** Nullable values with explicit handling
- **Function Types:** First-class functions for transformations

Term System

- **Literals:** Concrete values (strings, numbers, booleans)
- **Records:** Structured data with named fields
- **Lists and Collections:** Ordered and unordered data collections
- **Functions:** Transformation logic and computations

2.2.2. Schemas and Constraints

Input Schemas

Define the structure and constraints of your source data:

```
-- Example: CSV file structure
customerSchema = record "Customer" [
```

```
field "customerId" (required int32Type),
field "name" (required stringType),
field "email" (optional stringType),
field "registrationDate" (required stringType)
]
```

Output Schemas

Specify the target structure for transformed data:

```
-- Example: JSON API structure
customerApiSchema = record "CustomerAPI" [
  field "id" (required stringType),
  field "fullName" (required stringType),
  field "contactInfo" (optional contactInfoType),
  field "metadata" (required metadataType)
]
```

Validation Rules

Define business rules and constraints:

```
-- Example: Business validation rules
customerValidation = validationRules [
  rule "email_format" validateEmailFormat,
  rule "id_unique" validateUniqueId,
  rule "name_length" (validateLength 1 100)
]
```

2.2.3. Mappings and Transformations

Simple Field Mapping

Basic one-to-one field transformations:

```
customerMapping = mapping [
  mapField "customerId" "id" (transform toString),
  mapField "name" "fullName" identity,
  mapField "email" "contactInfo.email" identity
]
```

Complex Transformations

Advanced mapping logic with computations:

```
complexMapping = mapping [
```



```

computeField "fullName" $ \record ->
  concatenate [
    getField "firstName" record,
    " ",
    getField "lastName" record
  ],
computeField "ageCategory" $ \record ->
  let age = getField "age" record
  in if age < 18 then "minor" else "adult"
]

```

Conditional Logic

Handle different data scenarios:

```

conditionalMapping = mapping [
  conditionalField "status" $ \record ->
    case getField "isActive" record of
      True -> "active"
      False -> if hasField "suspensionReason" record
        then "suspended"
        else "inactive"
]

```

2.2.4. Data Quality and Validation

Input Validation

Ensure source data meets expectations:

```

inputValidation = validateInput [
  requireField "customerId",
  validateType "email" emailType,
  validateRange "age" 0 150,
  validatePattern "phone" phoneRegex
]

```

Output Validation

Verify transformed data correctness:

```

outputValidation = validateOutput [
  checkCompleteness requiredFields,
  checkConsistency relatedFields,
  checkBusinessRules businessConstraints
]

```

Error Handling

Define how to handle validation failures:

```
errorHandling = errorStrategy [  
  onMissingField logWarning,  
  onTypeError rejectRecord,  
  onBusinessRuleViolation flagForReview  
]
```

2.3. Development Environment and Tools

2.3.1. Setting Up Your Environment

Prerequisites

- Git for version control
- Pixi package manager (primary build system)
- Your chosen Hydra implementation (Haskell, Java, or Python)
- Text editor or IDE with appropriate language support

Initial Setup

```
# Clone the Hydra repository  
git clone https://github.com/CategoricalData/hydra  
cd hydra  
  
# Install all build dependencies  
pixi install  
  
# Validate the setup  
pixi run validate  
  
# Build all components  
pixi run build-all
```

Environment Selection

Choose your primary development environment:

```
# For complete functionality (Haskell)  
pixi install -e haskell  
pixi run -e haskell haskell-setup  
  
# For Java development  
pixi install -e java
```

```
# For Python development
pixi install -e python

# For Scala development
pixi install -e scala
```

2.3.2. Development Workflow

Typical Development Cycle

1. **Analysis:** Understand source and target data formats
2. **Schema Design:** Define input and output schemas
3. **Mapping Creation:** Implement transformation logic
4. **Testing:** Validate with sample data
5. **Optimization:** Improve performance and reliability
6. **Deployment:** Package for production use

Interactive Development

```
# Start interactive session for rapid prototyping
pixi run -e haskell stack-ghci

# In GHCi:
:l YourTransformation
testMapping sampleData
:reload -- after making changes
```

Testing and Validation

```
# Run transformation tests
pixi run -e haskell stack-test

# Test with sample data
pixi run -e haskell run-example-graphson

# Verify output quality
pixi run -e haskell check-copilot-data
```

2.3.3. Available Tools and Resources

Code Generation

```
# Generate language-specific code from schemas
```

```
pixi run -e haskell generate-java
pixi run -e haskell generate-haskell

# Generate documentation
pixi run -e haskell generate-docs
```

Data Processing

```
# Process sample datasets
pixi run -e haskell run-copilot-graphson
pixi run -e haskell run-example-graphson

# Validate existing data
pixi run -e haskell check-data
pixi run -e haskell check-graphson
```

Debugging and Exploration

```
# Explore available modules and functions
pixi run -e haskell haskell-modules
pixi run -e haskell haskell-functions

# Interactive debugging
pixi run -e haskell stack-ghci
-- :trace functionName
-- :browse ModuleName
```

2.4. Project Organization

2.4.1. Recommended Directory Structure

```
your-hydra-project/
├── data/
│   ├── input/           # Source data files
│   ├── samples/        # Sample datasets for testing
│   └── output/          # Generated output files
├── schemas/
│   ├── input/           # Input schema definitions
│   ├── output/          # Output schema definitions
│   └── validation/      # Validation rule definitions
├── transformations/
│   ├── mappings/        # Data mapping definitions
│   ├── functions/       # Custom transformation functions
│   └── pipelines/       # Complete transformation pipelines
├── tests/
│   └── unit/            # Unit tests for individual components
```

```

├── integration/ # End-to-end pipeline tests
├── data/        # Test data files
├── docs/
│   ├── schemas.md # Schema documentation
│   ├── mappings.md # Mapping documentation
│   └── deployment.md # Deployment instructions
├── scripts/
│   ├── build.sh # Build scripts
│   ├── test.sh # Testing scripts
│   └── deploy.sh # Deployment scripts
└── config/
    ├── development/ # Development configuration
    ├── staging/      # Staging configuration
    └── production/  # Production configuration

```

2.4.2. Configuration Management

Environment-Specific Configuration

```

# config/development/hydra.yaml
environment: development
input:
  source: "data/input/"
  format: "csv"
  validation: strict
output:
  destination: "data/output/"
  format: "json"
  compression: none
logging:
  level: debug
  output: console
performance:
  batch_size: 1000
  parallel_workers: 2

```

Production Configuration

```

# config/production/hydra.yaml
environment: production
input:
  source: "s3://data-bucket/input/"
  format: "avro"
  validation: strict
output:
  destination: "s3://data-bucket/output/"
  format: "parquet"
  compression: snappy

```

```
logging:
  level: info
  output: "logs/hydra.log"
performance:
  batch_size: 10000
  parallel_workers: 8
monitoring:
  metrics_endpoint: "http://metrics:9090"
  health_check_interval: 30
```

2.5. Best Practices

2.5.1. Schema Design

Keep Schemas Simple and Focused

- Use clear, descriptive field names
- Avoid deeply nested structures when possible
- Separate concerns into different schema modules
- Document business rules and constraints

Version Your Schemas

- Use semantic versioning for schema changes
- Maintain backward compatibility when possible
- Provide migration paths for breaking changes
- Test schema evolution scenarios

Validate Early and Often

- Validate input data before processing
- Check intermediate results during transformation
- Verify output data meets requirements
- Log validation failures with context

2.5.2. Transformation Design

Make Mappings Explicit

- Document the business logic behind each mapping
- Use descriptive names for transformation functions
- Avoid implicit conversions when possible
- Test edge cases and error conditions

Handle Errors Gracefully

- Define clear error handling strategies
- Provide meaningful error messages
- Implement fallback behaviors where appropriate
- Monitor and alert on error rates

Optimize for Maintainability

- Break complex transformations into smaller functions
- Use type signatures to document interfaces
- Write comprehensive tests for transformation logic
- Keep business logic separate from technical concerns

2.5.3. Performance Considerations

Design for Scale

- Consider memory usage for large datasets
- Implement streaming processing when possible
- Use appropriate data structures for your use case
- Profile and benchmark critical paths

Monitor and Measure

- Track processing times and throughput
- Monitor memory usage patterns
- Set up alerts for performance degradation
- Collect metrics on data quality issues

2.6. Common Patterns and Anti-Patterns

2.6.1. Effective Patterns

Schema-First Development

Start with clear schema definitions before implementing transformations.

Incremental Processing

Process data in manageable chunks to control memory usage.

Separation of Concerns

Keep parsing, transformation, and output generation as separate phases.

Comprehensive Testing

Test with realistic data volumes and edge cases.

2.6.2. Anti-Patterns to Avoid

Hardcoded Assumptions

Don't assume data will always match expected formats or ranges.

Monolithic Transformations

Avoid creating single, complex transformation functions.

Ignoring Data Quality

Don't skip validation and error handling.

Premature Optimization

Focus on correctness first, then optimize based on actual performance needs.

2.7. Next Steps

Once you understand these fundamentals:

1. **Practice with Examples:** Work through the provided transformation examples
2. **Start Small:** Begin with simple data format conversions
3. **Build Incrementally:** Add complexity gradually as you gain experience
4. **Engage with Community:** Join Discord discussions and ask questions
5. **Contribute Back:** Share your transformations and improvements

2.8. Related Documentation

- [Operator Examples](#) - Practical transformation examples
- [Data Transformation Patterns](#) - Common patterns and solutions
- [Performance Optimization](#) - Scaling and optimization strategies
- [Complete Operator Guide](#) - Comprehensive usage guide

Chapter 3. Part II: Getting Started with Data Transformation

Chapter 4. Transformation Examples

This document provides practical, step-by-step examples of data transformations using Hydra, from simple format conversions to complex schema mappings.

4.1. Your First Transformation: Relational to Graph

Let's walk through a practical example of transforming relational data into graph form, inspired by a veterinary office scenario.

4.1.1. The Scenario

Imagine you have data stored in CSV files:

- `pets.csv`: Information about pets (id, name, animal_type, owner_id)
- `owners.csv`: Information about owners (id, name, phone, afraid_of)
- `appointments.csv`: Scheduled appointments (id, pet_id, date, notes)
- `pet_friendships.csv`: Friendly relationships (pet_id_1, pet_id_2)
- `pet_fights.csv`: Conflict relationships (pet_id_1, pet_id_2)

You want to create a graph where:

- Owners, Pets, and Appointments become vertices
- Friendships and Fights become edges
- You can easily query relationships for scheduling purposes

4.1.2. Step 1: Define Your Input Schema

Create `InputSchema.hs` to describe your CSV structure:

```
module InputSchema where

import Hydra.Dsl.Types as Types
import Hydra.Core

-- Define the structure of each CSV file
petTableType :: Type m
petTableType = tableType [
    columnType "pet_id" integerType,
    columnType "name" stringType,
    columnType "animal_type" stringType,
    columnType "owner_id" integerType
]

ownerTableType :: Type m
ownerTableType = tableType [
```

```

columnType "owner_id" integerType,
columnType "name" stringType,
columnType "phone" stringType,
columnType "afraid_of" stringType
]

appointmentTableType :: Type m
appointmentTableType = tableType [
  columnType "appointment_id" integerType,
  columnType "pet_id" integerType,
  columnType "date" stringType,
  columnType "notes" stringType
]

friendshipTableType :: Type m
friendshipTableType = tableType [
  columnType "pet_id_1" integerType,
  columnType "pet_id_2" integerType
]

fightTableType :: Type m
fightTableType = tableType [
  columnType "pet_id_1" integerType,
  columnType "pet_id_2" integerType
]

```

4.1.3. Step 2: Define Your Output Schema

Create `OutputSchema.hs` to describe your target graph structure:

```

module OutputSchema where

import Hydra.Dsl.Types as Types
import Hydra.Ext.Pg.Model as Pg

-- Define the property graph schema
vetOfficeSchema :: Schema
vetOfficeSchema = schema "VetOffice" [
  -- Vertex types
  vertexType "Owner" [
    propertyType "name" stringType False,
    propertyType "phone" stringType False,
    propertyType "afraid_of" stringType True
  ],

  vertexType "Pet" [
    propertyType "name" stringType False,
    propertyType "animal_type" stringType False
  ],

```

```

vertexType "Appointment" [
  propertyType "date" stringType False,
  propertyType "notes" stringType True
],

-- Edge types
edgeType "owns" "Owner" "Pet" [],
edgeType "friendsWith" "Pet" "Pet" [],
edgeType "foughtWith" "Pet" "Pet" [],
edgeType "hasAppointment" "Pet" "Appointment" []
]

```

4.1.4. Step 3: Define Your Mapping

Create `Mapping.hs` to specify how data flows from input to output:

```

module Mapping where

import Hydra.Dsl.Terms as Terms
import qualified InputSchema as Input
import qualified OutputSchema as Output

-- Create vertices from CSV rows
createOwnerVertex :: Term -> Term
createOwnerVertex row = vertex "Owner" (stringLiteral ("owner_" ++ show (project
"owner_id" row))) [
  property "name" (project "name" row),
  property "phone" (project "phone" row),
  property "afraid_of" (project "afraid_of" row)
]

createPetVertex :: Term -> Term
createPetVertex row = vertex "Pet" (stringLiteral ("pet_" ++ show (project "pet_id"
row))) [
  property "name" (project "name" row),
  property "animal_type" (project "animal_type" row)
]

createAppointmentVertex :: Term -> Term
createAppointmentVertex row = vertex "Appointment" (stringLiteral ("appointment_" ++
show (project "appointment_id" row))) [
  property "date" (project "date" row),
  property "notes" (project "notes" row)
]

-- Create edges from relationship data
createOwnsEdge :: Term -> Term -> Term
createOwnsEdge petRow ownerRows =
  let ownerId = project "owner_id" petRow
      petId = project "pet_id" petRow

```

```

in edge "owns"
  (stringLiteral ("owner_" ++ show ownerId))
  (stringLiteral ("pet_" ++ show petId))
  []

createFriendshipEdge :: Term -> Term
createFriendshipEdge row =
  let pet1 = project "pet_id_1" row
      pet2 = project "pet_id_2" row
  in edge "friendsWith"
    (stringLiteral ("pet_" ++ show pet1))
    (stringLiteral ("pet_" ++ show pet2))
    []

-- Complete mapping specification
vetOfficeMapping :: Mapping
vetOfficeMapping = mapping [
  -- Map table rows to vertices
  mapTable "pets.csv" Input.petTableType createPetVertex,
  mapTable "owners.csv" Input.ownerTableType createOwnerVertex,
  mapTable "appointments.csv" Input.appointmentTableType createAppointmentVertex,

  -- Map relationship tables to edges
  mapTable "pet_friendships.csv" Input.friendshipTableType createFriendshipEdge,
  mapTable "pet_fights.csv" Input.fightTableType createFightEdge,

  -- Create ownership edges by joining pet and owner data
  joinTables "pets.csv" "owners.csv" "owner_id" createOwnsEdge
]

```

4.1.5. Step 4: Generate Your Graph

Use Hydra's built-in transformation capabilities to execute the mapping:

```

import Hydra.Ext.Demos.GenPG.Demo

main :: IO ()
main = do
  generateGraphSON
    "data/input"          -- Input directory
    vetOfficeMapping      -- Your mapping
    Output.vetOfficeSchema -- Your output schema
    "output/vet_office.json" -- Output file

```

This generates a GraphSON file that you can import into graph databases or visualization tools.

4.2. Advanced Transformation Patterns

4.2.1. Working with Complex Data Types

Nested JSON Structures

When working with nested JSON data, you can define recursive type structures:

```
-- Define a complex JSON structure
customerOrderSchema :: Type m
customerOrderSchema = recordType [
    field "customer_id" integerType,
    field "customer_info" customerInfoType,
    field "orders" (listType orderType)
]

customerInfoType :: Type m
customerInfoType = recordType [
    field "name" stringType,
    field "email" stringType,
    field "address" addressType
]

orderType :: Type m
orderType = recordType [
    field "order_id" integerType,
    field "items" (listType itemType),
    field "total" floatType
]
```


Handling Optional Fields

Use Hydra's optional type support for fields that may be missing:

```
userProfileType :: Type m
userProfileType = recordType [
    field "user_id" integerType,
    field "username" stringType,
    field "email" (optionalType stringType), -- May be null
    field "profile_picture" (optionalType stringType),
    field "preferences" (optionalType preferencesType)
]
```

4.2.2. Multi-Step Transformations

For complex transformations, you can chain multiple mapping steps:



```

-- First transformation: Raw data to intermediate format
step1Mapping :: Mapping
step1Mapping = mapping [
    mapTable "raw_events.csv" rawEventType normalizeEvent,
    mapTable "raw_users.csv" rawUserType normalizeUser
]

-- Second transformation: Intermediate format to final graph
step2Mapping :: Mapping
step2Mapping = mapping [
    mapData normalizedEvents createEventVertices,
    mapData normalizedUsers createUserVertices,
    createEdges "interactions" linkUsersToEvents
]

-- Execute the pipeline
transformationPipeline :: IO ()
transformationPipeline = do
    intermediate <- executeMapping step1Mapping "data/raw"
    final <- executeMapping step2Mapping intermediate
    exportGraph final "output/final_graph.json"
  
```

4.2.3. Schema Evolution and Migration

Handle evolving data schemas with version-aware mappings:

```
-- Version-aware schema mapping
migrateUserSchema :: SchemaVersion -> Term -> Term
migrateUserSchema version userData = case version of
  V1 -> migrationV1toV2 userData
  V2 -> migrationV2toV3 userData
  V3 -> userData -- Current version

migrationV1toV2 :: Term -> Term
migrationV1toV2 oldData = record [
  -- Preserve existing fields
  copyField "id" oldData,
  copyField "name" oldData,
  -- Add new required fields with defaults
  newField "email" (stringLiteral ""),
  newField "created_at" (stringLiteral "1970-01-01")
]

migrationV2toV3 :: Term -> Term
migrationV2toV3 oldData = record [
  -- Restructure data
  copyField "id" oldData,
  copyField "name" oldData,
  -- Combine fields
  newField "contact_info" $ record [
    field "email" (getField "email" oldData),
    field "phone" (getField "phone" oldData)
  ]
]
```

4.3. Working with Different Data Formats

4.3.1. CSV and Tabular Data

Reading CSV with Custom Parsers

```
-- Custom CSV parsing configuration
csvConfig :: CSVConfig
csvConfig = CSVConfig {
  csvDelimiter = ',',
  csvQuoteChar = '"',
  csvEscapeChar = '\\',
  csvHeaderRow = True,
  csvSkipRows = 0
}
```



```

-- Parse CSV with type validation
parseTypedCSV :: FilePath -> Type -> IO [Term]
parseTypedCSV path schema = do
  rawData <- parseCSVFile csvConfig path
  mapM (validateAndConvert schema) rawData

-- Custom field transformations
customerCSVMapping :: Mapping
customerCSVMapping = mapping [
  mapField "customer_id" "id" (transform stringToInt),
  mapField "full_name" "name" (transform titleCase),
  mapField "email_address" "email" (transform toLowerCase),
  computeField "display_name" $ \row ->
    concatenate [
      getField "first_name" row,
      " ",
      getField "last_name" row
    ]
]

```

4.3.2. JSON and Semi-Structured Data

Working with Variable JSON Schemas

```

-- Handle polymorphic JSON structures
parsePolymorphicJSON :: Value -> Term
parsePolymorphicJSON jsonValue =
  case getField "type" jsonValue of
    Just "user" -> parseUserObject jsonValue
    Just "organization" -> parseOrgObject jsonValue
    Just "event" -> parseEventObject jsonValue
    _ -> parseGenericObject jsonValue

-- Dynamic field extraction
extractDynamicFields :: Value -> [Term]
extractDynamicFields jsonValue =
  let baseFields = extractKnownFields jsonValue
      customFields = extractCustomFields jsonValue
  in baseFields ++ customFields

-- Schema inference from examples
inferJSONSchema :: [Value] -> Type
inferJSONSchema examples =
  let fieldTypes = analyzeFieldTypes examples
      optionalFields = identifyOptionalFields examples
  in buildRecordType fieldTypes optionalFields

```

4.3.3. XML and Hierarchical Data

Mapping XML Elements

```
-- XML to Hydra term conversion
parseXMLDocument :: XMLDocument -> Term
parseXMLDocument doc =
  let rootElement = xmlDocRoot doc
  in parseXMLElement rootElement

parseXMLElement :: XMLElement -> Term
parseXMLElement element = record [
  field "tag" (stringLiteral (xmlElementName element)),
  field "attributes" (parseAttributes (xmlElementAttrs element)),
  field "children" (list (map parseXMLNode (xmlElementChildren element)))
]

-- Handle mixed content
parseXMLNode :: XMLNode -> Term
parseXMLNode node = case node of
  XMLElement elem -> parseXMLElement elem
  XMLText text -> record [field "type" (stringLiteral "text"), field "content"
    (stringLiteral text)]
  XMLComment comment -> record [field "type" (stringLiteral "comment"), field
    "content" (stringLiteral comment)]

-- XML namespace handling
resolveXMLNamespace :: NamespaceMap -> QualifiedName -> String
resolveXMLNamespace nsMap qname =
  case lookupNamespace (qnameNamespace qname) nsMap of
    Just uri -> uri ++ "#" ++ qnameLocalName qname
    Nothing -> qnameLocalName qname
```

4.4. Message Translation and Integration

4.4.1. Building Message Translators

Protocol Buffer to JSON Translation

```
-- Define Protocol Buffer schema in Hydra
protobufPersonSchema :: Type
protobufPersonSchema = recordType [
  field "id" int32Type,
  field "name" stringType,
  field "email" stringType,
  field "phones" (listType phoneNumberType)
]
```

```

phoneNumberType :: Type
phoneNumberType = recordType [
    field "number" stringType,
    field "type" phoneTypeEnum
]

-- Define JSON output schema
jsonPersonSchema :: Type
jsonPersonSchema = recordType [
    field "id" stringType, -- Convert to string
    field "full_name" stringType,
    field "contact" contactInfoType
]

contactInfoType :: Type
contactInfoType = recordType [
    field "email" stringType,
    field "phone_numbers" (listType stringType)
]

-- Translation mapping
protobufToJsonMapping :: Mapping
protobufToJsonMapping = mapping [
    -- Basic field mappings with transformations
    mapField "id" "id" (transform intToString),
    mapField "name" "full_name" identity,

    -- Complex nested mapping
    computeField "contact" $ \protobufRecord -> record [
        field "email" (getField "email" protobufRecord),
        field "phone_numbers" (transformPhoneList (getField "phones" protobufRecord))
    ]
]

transformPhoneList :: Term -> Term
transformPhoneList phoneList =
    list $ map extractPhoneNumber (getListElements phoneList)

extractPhoneNumber :: Term -> Term
extractPhoneNumber phoneRecord = getField "number" phoneRecord

```

4.4.2. Event Stream Processing

Real-Time Data Transformation

```

-- Stream processing pipeline
eventStreamProcessor :: EventStream -> TransformationPipeline -> IO ()
eventStreamProcessor stream pipeline = do
    processStream stream $ \event -> do

```

```

-- Parse incoming event
parsedEvent <- parseEvent event

-- Apply transformations
transformedEvent <- applyPipeline pipeline parsedEvent

-- Validate output
validatedEvent <- validateEvent transformedEvent

-- Send to downstream systems
publishEvent validatedEvent

-- Event schema with polymorphic handling
eventSchema :: Type
eventSchema = unionType "Event" [
  variant "UserAction" userActionType,
  variant "SystemEvent" systemEventType,
  variant "ErrorEvent" errorEventType
]

-- Conditional processing based on event type
processEventByType :: Term -> IO Term
processEventByType event = case getVariant event of
  ("UserAction", userData) -> processUserAction userData
  ("SystemEvent", sysData) -> processSystemEvent sysData
  ("ErrorEvent", errData) -> processErrorEvent errData
  (unknownType, _) -> throwError $ UnknownEventType unknownType

-- Batch processing for efficiency
batchProcessor :: Int -> [Event] -> IO [TransformedEvent]
batchProcessor batchSize events = do
  let batches = chunksOf batchSize events
  results <- mapConcurrently processBatch batches
  return $ concat results

```

4.4.3. Data Validation and Quality

Schema Validation

```

-- Comprehensive validation rules
validationRules :: [ValidationRule]
validationRules = [
  -- Type validation
  validateType "email" emailType,
  validateType "phone" phoneType,
  validateType "date" iso8601DateType,

  -- Range validation
  validateRange "age" 0 150,

```

```

validateRange "score" 0.0 100.0,

-- Pattern validation
validatePattern "zipcode" zipCodeRegex,
validatePattern "ssn" ssnRegex,

-- Cross-field validation
validateRelationship "start_date" "end_date" datePrecedence,
validateRelationship "min_value" "max_value" numericPrecedence,

-- Business rules
validateBusinessRule "customer_age_product" ageProductCompatibility,
validateBusinessRule "region_shipping" regionShippingRules
]

-- Custom validation functions
emailType :: ValidationFunction
emailType value = case value of
  StringLiteral email -> validateEmailFormat email
  _ -> ValidationError "Email must be a string"

validateEmailFormat :: String -> ValidationResult
validateEmailFormat email =
  if email =~ emailRegex
  then ValidationSuccess
  else ValidationError "Invalid email format"

-- Validation pipeline
validateDataPipeline :: ValidationRules -> Term -> Either [ValidationError] Term
validateDataPipeline rules inputData = do
  -- Apply all validation rules
  results <- mapM (applyRule inputData) rules

  -- Collect errors
  let errors = [err | ValidationError err <- results]

  if null errors
  then Right inputData
  else Left errors

```

4.5. Performance and Optimization

4.5.1. Streaming and Large Datasets

```

-- Streaming transformation for large files
streamingTransformation :: FilePath -> FilePath -> Mapping -> IO ()
streamingTransformation inputPath outputPath mapping = do
  withFile inputPath ReadMode $ \inputHandle ->
    withFile outputPath WriteMode $ \outputHandle -> do

```

```

    sourceHandle inputHandle
    $$ transformConduit mapping
    $$ sinkHandle outputHandle

-- Memory-efficient batch processing
batchTransform :: Int -> [Term] -> Mapping -> [Term]
batchTransform batchSize items mapping =
    let batches = chunksOf batchSize items
    in concatMap (applyMapping mapping) batches

-- Parallel processing
parallelTransform :: Mapping -> [Term] -> IO [Term]
parallelTransform mapping items = do
    let chunks = chunksOf 1000 items
    results <- mapConcurrently (return . applyMapping mapping) chunks
    return $ concat results

-- Progress monitoring
transformWithProgress :: Mapping -> [Term] -> IO [Term]
transformWithProgress mapping items = do
    progressRef <- newIORef 0
    total <- return $ length items

    results <- forM items $ \item -> do
        result <- return $ applyMappingSingle mapping item
        modifyIORef progressRef (+1)
        progress <- readIORef progressRef
        when (progress `mod` 100 == 0) $
            putStrLn $ "Processed " ++ show progress ++ "/" ++ show total
        return result

    return results

```

4.5.2. Caching and Memoization

```

-- Memoized transformations for expensive operations
memoizedTransform :: (Term -> Term) -> IO (Term -> IO Term)
memoizedTransform f = do
    cache <- newIORef Map.empty
    return $ \input -> do
        cacheMap <- readIORef cache
        case Map.lookup input cacheMap of
            Just result -> return result
            Nothing -> do
                let result = f input
                modifyIORef cache (Map.insert input result)
                return result

-- LRU cache for limited memory usage

```

```
createLRUCache :: Int -> IO (LRUCache Term Term)
createLRUCache maxSize = newLRUCache maxSize

-- Cached schema validation
cachedValidation :: LRUCache Type ValidationFunction -> Type -> Term -> IO
ValidationResult
cachedValidation cache schema term = do
  validator <- lookupLRU cache schema $ compileValidator schema
  return $ validator term
```

4.6. Related Documentation

- [Operator Fundamentals](#) - Core concepts and setup
- [Data Transformation Patterns](#) - Common patterns and solutions
- [Performance Optimization](#) - Scaling and optimization strategies
- [Complete Operator Guide](#) - Full operator documentation

Chapter 5. Part III: Development Workflow with Pixi Tasks

5.1. Working with Generated Code

For complex transformations, you'll often need to generate code from Hydra definitions:

```
# Generate Java code from Hydra definitions
pixi run -e haskell generate-java

# Generate Haskell code from Hydra definitions
pixi run -e haskell generate-haskell

# Start interactive Haskell session for development
pixi run -e haskell stack-ghci
```

5.2. Data Processing Tasks

```
# Check available demo data
pixi run -e haskell check-data

# Generate GraphSON sample data
pixi run -e haskell run-example-graphson
pixi run -e haskell run-copilot-graphson

# Verify generated output
pixi run -e haskell check-graphson
pixi run -e haskell check-copilot-data
pixi run -e haskell check-sales-data
```

5.3. Testing Your Transformations

```
# Run Haskell tests to validate transformations
pixi run -e haskell stack-test
pixi run -e haskell hydra-test
pixi run -e haskell ext-test

# Test individual transformation modules
pixi run -e haskell stack-ghci
-- In GHCi: :l YourTransformation && runTests
```


5.4. Exploring Hydra Modules and Functions

```
# List available Haskell modules
pixi run -e haskell haskell-modules

# Find functions in generated code
pixi run -e haskell haskell-functions

# Browse modules interactively
pixi run -e haskell stack-ghci
-- In GHCi: :browse Hydra.Core.Model
```

5.5. Building and Packaging

```
# Build all components for your target platform
pixi run build-all

# Build specific components
pixi run build-java      # For Java-based deployments
pixi run build-python    # For Python-based deployments
pixi run build-haskell   # For Haskell-based deployments

# Clean build artifacts
pixi run clean
```

5.6. Advanced Workflow Patterns

5.6.1. Iterative Development and Testing

```
# Clean and rebuild cycle for development
pixi run clean
pixi run build-all

# Interactive development with immediate feedback
pixi run -e haskell stack-ghci
-- In GHCi: :reload after making changes

# Test specific transformations
pixi run -e haskell hydra-test
pixi run -e haskell ext-test
```

5.6.2. Working with Large Datasets

```
# Check available sample data
```

```
pixi run -e haskell check-data
pixi run -e haskell check-copilot-data
pixi run -e haskell check-sales-data

# Generate test datasets
pixi run -e haskell run-copilot-graphson
pixi run -e haskell run-example-graphson

# Verify output quality
pixi run -e haskell check-graphson
```

5.6.3. Debugging Transformations

```
# Start interactive session for debugging
pixi run -e haskell stack-ghci

# In GHCi - load your transformation module
:l YourTransformationModule
:browse YourTransformationModule
:type yourTransformationFunction

# Clear caches if encountering issues
pixi run -e haskell clear-stack-cache-all
pixi run -e haskell stack-clean
```

5.6.4. Performance Optimization

```
# Build with optimizations
pixi run -e haskell stack-build

# Profile memory usage (in GHCi)
pixi run -e haskell stack-ghci
-- :set +s -- Show timing/memory stats
-- :set +t -- Show types

# Clean previous builds for fresh benchmarks
pixi run -e haskell stack-clean
pixi run build-all
```

Chapter 6. Part IV: Production Deployment

6.1. Packaging for Production

6.1.1. Creating Distributable Packages

```
# Build production-ready packages
pixi run build-all

# Create platform-specific distributions
pixi run -e java gradle build
pixi run -e haskell stack install

# Verify package integrity
pixi run validate
```

6.1.2. Configuration Management

- **Environment-specific configurations:** Development, staging, production
- **Schema versioning:** Handle evolving data schemas
- **Error handling strategies:** Define how to handle validation failures and data quality issues
- **Performance tuning:** Optimize for your specific workloads and data volumes

6.2. Monitoring and Maintenance

6.2.1. Health Monitoring

- **Data quality metrics:** Track validation failures, schema violations
- **Performance monitoring:** Processing times, throughput, memory usage
- **Error tracking:** Log and alert on transformation failures
- **Schema drift detection:** Monitor for unexpected changes in input data

6.2.2. Maintenance Tasks

```
# Regular maintenance tasks
pixi run clean                # Clean build artifacts
pixi run validate             # Validate system health
pixi run -e haskell stack-test # Run regression tests

# Update dependencies
pixi install                  # Update pixi dependencies
pixi run -e haskell stack update # Update Haskell dependencies
```

Chapter 7. Appendices

Appendix A: Operator Resources

- **Main Repository:** <https://github.com/CategoricalData/hydra>
- **LambdaGraph Discord:** <https://bit.ly/lg-discord>
- **Community Examples:** Available in the hydra-ext subproject
- **G.V() Visualization:** <https://gdotv.com> for graph visualization
- **Pixi Documentation:** <https://pixi.sh> for build system reference

Appendix B: Common Transformation Patterns

7.B.1. Data Format Conversions

- CSV to JSON transformation patterns
- XML to Protocol Buffer mappings
- Relational to graph data migrations
- API format translations

7.B.2. Schema Evolution Patterns

- Forward compatibility strategies
- Backward compatibility maintenance
- Version migration workflows
- Data validation and quality assurance

7.B.3. Performance Optimization Patterns

- Streaming data processing
- Batch optimization techniques
- Memory management strategies
- Parallel processing approaches

Appendix C: Troubleshooting Guide

7.C.1. Common Issues

- **Build failures:** Check dependencies and environment setup
- **Memory issues:** Use streaming processing for large datasets
- **Performance problems:** Profile and optimize critical paths

- **Schema mismatches:** Validate input data against expected schemas
- **Type errors:** Ensure proper type definitions and mappings

7.C.2. Getting Help

- Join the Discord community for real-time assistance
- Check existing issues in the GitHub repository
- Review the comprehensive developer documentation
- Examine working examples in the hydra-ext project

Colophon

This operator guide covers the practical application of Hydra for data transformation tasks. It focuses on real-world scenarios and provides step-by-step guidance for operators using existing Hydra implementations.

For the latest version of this guide and related documentation, see the [docs/](#) directory in the Hydra repository.